

Marshall Zhang (He/Him/His)

Email: marshall.yang.zhang@outlook.com *Website:* marshallzzz.com

Looking for opportunities in AI Agent and platform product management

Professional Summary

Product Manager with experience building and scaling execution platforms for distributed data systems, including scheduling, observability, and execution identity in Microsoft Fabric. Proven track record of transforming fragmented system experiences into unified, high-adoption platforms (400K+ MAU), while enabling enterprise-scale workloads through thoughtful tradeoffs across performance, reliability, and usability.

Strong focus on applying AI to real-world scenarios, including workflow automation, intelligent monitoring, and AI-assisted product development, with an emphasis on identifying high-impact use cases and integrating AI into end-to-end systems rather than isolated features.

Experience

Microsoft — Beijing, China

Program Manager, C+AI, Microsoft Fabric

Mar 2022 – Present

Fabric Job Platform (Monitoring Hub)

- Identified fragmented observability across Fabric workloads as a key barrier to efficient debugging, and drove creation of Monitoring Hub as a unified platform replacing multiple workload-specific monitoring experiences
- Scaled adoption to 400K+ MAU (~80% of Fabric MAU) by prioritizing core user scenarios (job discovery, failure diagnosis, cross-workload debugging) over workload-specific feature parity
- Eliminated multi-tool troubleshooting workflows by consolidating execution signals across all executable workloads, enabling monitoring and investigation from a single entry point
- Evolved observability strategy beyond raw signal aggregation by introducing AI-driven monitoring scenarios (failure summarization, pattern clustering, anomaly detection), reducing manual effort in diagnosing recurring issues
- Prioritized proactive visibility into system health by building notification and alerting mechanisms, enabling earlier detection of failures and reducing reliance on manual monitoring

Job Scheduling & Execution Scalability

- Enabled support for large-scale enterprise pipelines (1000+ executions per minute per user, long-running jobs) by redefining concurrency and throttling models to handle high-throughput, multi-workload scenarios
- Unblocked ADF / HDInsight → Fabric migration at scale by identifying and resolving key execution constraints (concurrency limits, API throttling, reliability gaps) that prevented customers from running production workloads
- Redefined API rate limiting strategy from user-based to item- and workspace-level controls, balancing user experience with platform protection under heavy and potentially abusive workloads
- Prioritized execution reliability for enterprise automation scenarios by addressing failure patterns such as credential expiration and long-running job lifecycle issues

AI-Driven Workflow & Productivity

- Identified signal overload in customer feedback channels (~100 posts/week, ~20% actionable) and drove AI-assisted triage workflow using keyword filtering and semantic analysis, reducing manual review volume to ~20 posts/week (~80% noise reduction) and shrinking response latency from days to hours (~10× improvement)
- Recognized opportunity to improve development workflow efficiency and enabled an AI-assisted product delivery pipeline spanning PM, design, and engineering, automating stages from requirement synthesis to spec generation, design prototyping, and implementation support
- Reduced spec authoring cycle (including review) from ~1 week to ~1 day (~5× faster) and design iteration from ~1–4 weeks to ~1–3 days by integrating AI into well-defined and repeatable workflow stages
- Defined adoption boundaries for AI in product development, where ~10% of well-scoped work items can be fully automated while all items leverage AI assistance at some stage, shifting PM focus toward scenario definition, prioritization, and outcome evaluation
- Contributed to Fabric unified Copilot initiative by aligning execution platform capabilities with a cross-workload, context-aware AI assistant, supporting conversational workflows and persistent context experiences

Microsoft — Beijing, China

Software Engineer (Frontend), Fabric

Mar 2021 – May 2022

- Developed UI experiences for Azure Fabric platform (Homepage, Playground), improving usability and interaction quality
- Built reusable frontend components and contributed to end-to-end testing, accessibility compliance, and unit testing practices

Microsoft — Beijing, China

Software Engineer (Frontend), Azure Streaming

Mar 2020 – May 2021

- Delivered UI integration for Azure Streaming Service within Azure Synapse
- Built DAG-like visualization experience to represent streaming job execution relationships

Microsoft — Redmond, WA, US

Software Engineer (Backend), HDInsight

Jan 2019 – Mar 2020

- Worked on Hadoop infrastructure, including logging systems and job queue management
- Supported open-source components and production systems; acted as on-call engineer (DRI)

Dataguise Inc. — Fremont, CA, US

Software Engineer

Jul 2015 – Dec 2018

- Developed Hadoop-based data security and compliance solutions (GDPR, sensitive data protection)
- Led Scrum processes, project tracking, and production support

Side Project

AI Job Search Copilot / Resume Assist

2026

- Built an AI job-search assistant that compares resume/profile context with up to 10 JDs, ranks roles by success probability, and generates per-JD resume improvement suggestions.
- Designed an agent-ready workflow with LLM structured extraction, RAG profile evidence, tool/function calling, trace logging, schema validation, and human approval for external actions.
- Deployed Resume Assist on marshallzzz.com/resumeAssist to demonstrate how existing job-search workflows can be AI-enabled while keeping validation, safety boundaries, and explainability.

Education

SUNY Binghamton University, Binghamton, NY, US

Master of EECE

GPA: 4.0 / 2013 – 2015

Beijing University of Posts and Telecommunications, Beijing, CN

BS of EE

GPA: 3.2 / 2009 – 2013